

SUJAI HIEMATH

sujai.github.io
sh2583@cornell.edu
[Google Scholar](#)
[LinkedIn](#)

TL;DR

- 3rd year ORIE PhD student at Cornell Tech, working on causality and LLMs.
- Published 3 first-author papers (NeurIPS 2024, UAI 2025, NeurIPS 2025) on sample-efficient and robust causal discovery methods within 1.5 years.
- Currently interning at Amazon Research Tübingen, where I released a preprint on LLM-aided causal discovery within 4 months.

PUBLICATIONS AND PREPRINTS

1. **Hiremath, S.***, et al. From Guess2Graph: When and How Can Unreliable Experts Safely Boost Causal Discovery in Finite Samples? *arXiv preprint*, 2025.
2. Meier, D.* and **Hiremath, S.***, et al. When Additive Noise Meets Unobserved Mediators: Bivariate Denoising Diffusion for Causal Discovery. *Thirty-Ninth Annual Conference on Neural Information Processing Systems*, 2025.
3. **Hiremath, S.***, et al. LoSAM: Local Search in Additive Noise Models with Mixed Mechanisms and General Noise for Global Causal Discovery. *Proceedings of the Forty-first Conference on Uncertainty in Artificial Intelligence*, 2025.
4. **Hiremath, S.***, et al. Hybrid Top-Down Global Causal Discovery with Local Search for Linear and Nonlinear Additive Noise Models. *Thirty-Eighth Annual Conference on Neural Information Processing Systems*, 2024.

WORK EXPERIENCE

- Applied Scientist Intern** | Amazon Research Tübingen, Germany 06.2025 - Present
- Managers: [Dr. Dominik Janzing](#), [Dr. Shiva Kasiviswanathan](#), [Dr. Elke Kirschbaum](#).
 - Developed a method leveraging LLMs as unreliable experts to improve causal discovery in finite samples; released preprint in <4 months.
 - Currently investigating causal reinforcement learning approaches for sample-efficient training of LLMs in low-data regimes.
- PhD Student Researcher** | Cornell Tech 11.2023 - Present
- PIs: [Dr. Kyra Gan](#), [Dr. Promit Ghosal](#).
 - Developed 3 novel causal discovery algorithms that use diffusion models, independence tests, and nonparametric regression techniques.
 - Published 3 first-author papers at NeurIPS (2024, 2025) and UAI (2025) on improving finite-sample causal structure learning while relaxing assumptions.

EDUCATION

- Cornell Tech | New York, NY** 2024 - 2028
PhD in Operations Research and Information Engineering | GPA: 3.9 (expected)
• Areas: AI, Causal Discovery, Causal Inference
- Cornell University | Ithaca, NY** 2023 - 2024
PhD in Operations Research and Information Engineering
- California Institute of Technology | Pasadena, CA** 2019 - 2023
BS in Applied and Computational Mathematics | GPA: 4.0
• Areas: Machine Learning, Mathematical Modelling, Deep Learning

SERVICE & AWARDS

Service: Reviewer for NeurIPS 2025, ICLR 2025, AISTATS 2025
Awards: NeurIPS Top Reviewer 2025 | Cornell Fellowship 2023 | Thomas J. Watson Fellowship, IBM (2019-2022) | SURF Fellowship, Caltech (2020, 2021)